
PostgreSQL Foreign Data Wrapper development with Python.

Toufee Ockards

@tockards (Twitter handle, but I mostly tweet music videos)

PostgreSQL



- The world's most advanced open source database.
-

Foreign Data Wrapper (FDW)

- SQL standard that deals with how a database management system can integrate data stored outside the database.
 - **SQL/MED**
-

Python



- The world's most advanced open source language.
-

Foreign/Federated Table

- **Storage engine** which allows a user to create a table that is a **local representation of a foreign (remote) table**
 - a transparent access method for external data
-

Data

- Flat Files e.g CSV files, Log Files, /etc/passwd_(kappa) 
 - RSS feeds, Emails, Websites, APIs
 - Directories, file shares
 - Other Databases, Services
-

You would like

```
postgres=# select some_data from external_data_source;
```

Multicorn FDW

- PostgreSQL Extension.
- Allows you to Write FDW's in python.



MULTICORN

You have

```
Year,Make,Model,Length  
1997,Ford,E350,2.34  
2000,Mercury,Cougar,2.38
```

- A csv file sitting in /tmp (because mongo)
-

CSV Example (Python)

```
from . import ForeignDataWrapper
from .utils import log_to_postgres
from logging import WARNING
import csv

class CsvFdw(ForeignDataWrapper):
    """A foreign data wrapper for accessing csv files.

    Valid options:
    - filename : full path to the csv file, which must be readable
      by the user running postgresql (usually postgres)
    - delimiter : the delimiter used between fields.
      Default: ","
    """

    def __init__(self, fdw_options, fdw_columns):
        super(CsvFdw, self).__init__(fdw_options, fdw_columns)
        self.filename = fdw_options["filename"]
        self.delimiter = fdw_options.get("delimiter", ",")
        self.quotechar = fdw_options.get("quotechar", "'")
        self.skip_header = int(fdw_options.get('skip_header', 0))
        self.columns = fdw_columns
```

CSV Example (Python)

```
def execute(self, quals, columns):
    with open(self.filename) as stream:
        reader = csv.reader(stream, delimiter=self.delimiter)
        count = 0
        checked = False
        for line in reader:
            if count >= self.skip_header:
                if not checked:
                    # On first iteration, check if the lines are of the
                    # appropriate length
                    checked = True
                    if len(line) > len(self.columns):
                        log_to_postgres("There are more columns than "
                                       "defined in the table", WARNING)
                    if len(line) < len(self.columns):
                        log_to_postgres("There are less columns than "
                                       "defined in the table", WARNING)
                yield line[:len(self.columns)]
            count += 1
```

CSV Example (SQL)

```
CREATE SERVER csv_srv foreign data wrapper multicorner options (  
    wrapper 'mycode.csvfdw.CsvFdw'  
);
```

CSV Example (SQL)

```
create foreign table csvtest (  
    year numeric,  
    make character varying,  
    model character varying,  
    length numeric  
) server csv_srv options (  
    filename '/tmp/test.csv',  
    skip_header '1',  
    delimiter ',');
```

CSV Example (SQL)

```
select * from csvtest;
```

year	make	model	length
1997	Ford	E350	2.34
2000	Mercury	Cougar	2.38

(2 lines)

SLOC Python 44
SQL 3 Statements

What else

- Use your ORM, BI Tool, Kitchen Sink above it.
 - PostgreSQL 9.3 +
 - Insert, Update, Delete
 - SQLAlchemy FDW
 - LDAP FDW
 - Docker FDW
 - PGOSQUERY ([like](#) facebook osquery_(see what I did with the like))
 - Redis-FDW
 - Mongo-FDW
 - https://wiki.postgresql.org/wiki/Foreign_data_wrappers
-



SKA AFRICA
SQUARE KILOMETRE ARRAY

SKA SA Control and Monitoring Use Case

- We currently archive our sensor data to HDF5 Files and have legacy archived files in CCSV Format
- Query these Files via a server and use a PostgreSQL FDW to represent it in a table.
- **Time Series Sensor Data.**
 - sensor_name, ***sample_ts**, values_ts, status, value



What we liked

- SQL interface remains the same.
- One interface for querying multiple data sources.
 - HDF5 File Format.
 - Inhouse CCSV Format
- The power of SQL.





What ~~we~~ I did not like.

- Writing Unit Tests for the FDW.





Future

- CEPH RADOS(distributed object store)
- SQL interface remains the same.





Requirements

- PostgreSQL 9.1+
- PostgreSQL Development Packages
- Python Development Packages
- Python 2.7 >= or Python 3.3 or default python





We are hiring.

- <http://www.ska.ac.za/vacancies/>